

Data Mining Assignment 2

Michael (吳忠賢)

Assignment 2: Cross-Validation, SVM, FP-Growth, K-Means + ARM

Student ID: 314540066

GitHub Link: <https://github.com/Ashurali/DM2026-Assignment-1-MKS>

Q1. K-fold Cross-Validation (15 points)

Notebook: Real_World_Classification.ipynb

Q1(a) 5-fold CV grid (4×4 $lr \times \lambda$, L2 regularisation)

CV is run on the training set only (test set held out), with `KFold(n_splits=5, shuffle=True, random_state=40)` and `n_iteration=10000` per fold. Mean accuracy across the 5 folds, by configuration:


$lr \setminus \lambda$	1	2	4	8
0.005	0.7171	0.7200	0.7257	0.7229
0.01	0.7229	0.7286	0.7257	0.7171
0.1	0.7229	0.7257	0.7257	0.7171
0.5	0.7229	0.6914	0.5343	0.4657

CV winner is $lr=0.01$, $\lambda=2.0$ (0.7286). Four configurations tie at 0.7257 ($lr=0.005, \lambda=4$, $lr=0.01, \lambda=4$, $lr=0.1, \lambda=2$, $lr=0.1, \lambda=4$); the notebook's top-2 argsort picks $lr=0.01$, $\lambda=4.0$ as Top 2. We evaluate the winner plus all four tied configurations on the held-out test set in Q1(b).

Q1(b) Top-2 + tied configs on the test set

Config	Train acc	Test acc	Test Precision	Test Recall	Test F1
Top 1: lr=0.01, $\lambda=2.0$	—	0.7600	0.7381	0.8158	0.7750
Top 2: lr=0.01, $\lambda=4.0$	—	0.7400	0.7176	0.8026	0.7578
Tied: lr=0.005, $\lambda=4.0$	—	0.7533	0.7241	0.8289	0.7730
Tied: lr=0.1, $\lambda=4.0$	—	0.7400	0.7176	0.8026	0.7578
Tied: lr=0.1, $\lambda=2.0$	—	0.7533	0.7294	0.8158	0.7702

Confusion-matrix charts (Top 1 on the left, best-recall tie on the right):



Q1(c) Observations

On the CV step. The 4-way tie at CV=0.7257 ($\lambda=4$ across all three smaller learning rates plus $\lambda=2$ at lr=0.1) initially suggested $\lambda=4$ gave a consistent boost. The CV winner lr=0.01, $\lambda=2.0$ only led the runners-up by 0.0029 —within CV noise. The most striking pattern is the **lr $\times$ λ interaction at lr=0.5**: the model converges normally at $\lambda=1$ (0.7229) but degrades at $\lambda=2$ (0.6914), and outright fails to converge at $\lambda=4$ (0.5343) and $\lambda=8$ (0.4657). At smaller learning rates (≤ 0.1) the same λ values are all fine. This is a step-size $\times$ regularisation effect —large lr combined with strong reg_lambda makes the gradient updates overshoot, not an over-regularisation problem at $\lambda=8$ *per se*.

On the CV $\rightarrow$ test transfer. Three observations: 1. The CV-picked Top 1 (lr=0.01, $\lambda=2.0$) is best on test too (Acc 0.76, F1 0.7750). CV did transfer for the very best. 2. The CV-picked Top 2 (lr=0.01, $\lambda=4.0$) is actually the **worst** of the tie on test (Acc 0.74). The “consistent $\lambda=4$ boost” suggested by CV was CV-noise, not generalisation. 3. lr=0.01, $\lambda=4.0$ and lr=0.1,

$\lambda=4.0$ produce **identical test metrics** —same converged minimum; learning rate only controls speed, not destination, given enough iterations.

Lr-vs-recall coupling. $\text{lr}=0.005$, $\lambda=4.0$ gives the highest recall (0.8289) of any configuration. A slower learning rate combined with stronger regularisation pushes the decision boundary to predict the positive class more readily. Accuracy moves little with lr; recall is the lr-sensitive metric here.

Q2. SVM on mobile_price.csv (15 points)

Notebook: Mobile_Price_SVM.ipynb

Data is shuffled with `random_state=42` and split 60/20/20 into train/val/test. All 20 features are kept; `price_range` is the 4-way target.

Q2(a) SVC(C=1.0) baseline

Split	Accuracy	Macro F1
Train	0.9842	0.9842
Validation	0.8700	0.8698
Test	0.8725	0.8727

Train-val gap of ~ 11.4 pts indicates moderate overfitting at $C=1.0$; val and test agree closely.

Q2(b) C sweep

$C \in \{0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000\}$. Results (5-fold CV on train + final test):

C	val_acc	CV mean	CV std	test_acc
0.001	0.7075	0.2544	0.0496	0.6625
0.01	0.7075	0.2544	0.0496	0.6625
0.1	0.7075	0.6113	0.0357	0.6625
1.0	0.8700	0.8538	0.0023	0.8725
10.0	0.8500	0.8619	0.0252	0.8725
100.0	0.8500	0.8612	0.0247	0.8725
1000+	0.8500	0.8612	0.0247	0.8725





Figure 1: C sweep

Q2(c) Final pick

$C = 1.0$ by the TA-suggested procedure (highest val_acc), $C = 10$ by 5-fold CV (highest CV mean). Both deliver identical test accuracy (0.8725) —at $C \geq 1$ the SVM has saturated the linearly-separable signal in the data. The conservative choice is $C = 1.0$ because (a) it yields the smallest train-val gap, (b) at higher C the CV std jumps to 0.025 (10 $\times$ higher than at $C=1$), so $C=10$ wins CV mean only by being lucky on average.



Q3. Association Rule Mining (15 points)

Notebook: Mobile_Price_FPgrowth.ipynb

Subset to `price_range == 1` (500 samples). Discretise `{ram, int_memory, px_width, battery_power}` into low/medium/high using a 3:4:3 split of the value range. Build per-row transactions, then run `mlxtend.frequent_patterns.fpgrowth`.

Q3(a) Frequent patterns (support ≥ 0.3)

8 itemsets:

size	support	items
1	0.682	ram_medium
1	0.416	px_width_medium
1	0.414	battery_power_medium
1	0.412	int_memory_medium
2	0.318	battery_power_medium, ram_medium
1	0.316	int_memory_low
1	0.308	battery_power_low
2	0.306	px_width_medium, ram_medium

Q3(b) Association rules (support ≥ 0.3 , confidence ≥ 0.4 , lift ≥ 0.8)

4 rules:

Antecedents	Consequents	support	confidence	lift
battery_power_medium	ram_medium	0.318	0.7681	1.1263
ram_medium	battery_power_medium	0.318	0.4663	1.1263
px_width_medium	ram_medium	0.306	0.7356	1.0786
ram_medium	px_width_medium	0.306	0.4487	1.0786

Q3(c) Observations

In `price_range == 1` (mid-priced phones), `ram_medium` is dominant —68.2% of these phones have RAM in the middle bucket. The rules show `battery_power_medium → ram_medium` and `px_width_medium → ram_medium` both with confidence ≈ 0.75 and lift > 1 , meaning these mid-tier feature levels co-occur with mid-RAM at slightly above chance. Lift is only 1.08–1.13 —not a strong association, just a mild positive co-tendency. The rules' reverse direction (`ram_medium → ...`) has confidence around 0.45, which makes sense: most mid-priced phones do have medium RAM, but only some have medium battery/px_width.

The `ram_medium` saturation in this class will turn out to matter for Q5: it overlaps with `ram_medium` regions of class 2, which is the source of the K-Means ceiling we eventually overcome.

Q4. PCA + K-Means (20 points)

Notebook: `Mobile_Price_PCA_KMeans.ipynb`

Q4(a) Standardise features (z-score)

Split `mobile_price.csv` into the 20 features `X` and the `price_range` label `y`. Standardise `X` with `sklearn.preprocessing.StandardScaler` so every feature has mean 0 and unit variance.

Raw features (first 5 rows, first 6 columns):

	battery_power	blue	clock_speed	dual_sim	fc	four_g
0	842	0	2.2	0	1	0
1	1021	1	0.5	1	0	1
2	563	1	0.5	1	2	1
3	615	1	2.5	0	0	0
4	1821	1	1.2	0	13	1

After z-score standardisation (X_{std} , same rows/columns):

	battery_power	blue	clock_speed	dual_sim	fc	four_g
0	-0.9026	-0.9900	0.8308	-1.0192	-0.7625	-1.0440
1	-0.4951	1.0101	-1.2531	0.9812	-0.9929	0.9579
2	-1.5377	1.0101	-1.2531	0.9812	-0.5321	0.9579
3	-1.4193	1.0101	1.1985	-1.0192	-0.9929	-1.0440
4	1.3259	1.0101	-0.3950	-1.0192	2.0023	0.9579

Per-column verification: every standardised column has mean ≈ 0 and std = 1.0, confirming the z-score transform succeeded. Standardisation is required before PCA so that high-variance raw features (e.g. ram in MB, px_height in pixels with values up to thousands) don't dominate the principal components purely because of their unit scale.

Q4(b) PCA scatter coloured by class

Fit $\text{PCA}(n_components=2)$ on the standardised features and project all 2000 samples onto (PC1, PC2). Colour each point by its true price_range label and add a legend.




Figure 2: PCA scatter coloured by true class

PCA-2 explained variance: [0.0839, 0.0811], cumulative = 0.1650

Only **16.5%** of variance lives in the top 2 components, with PC1 = 8.4% and PC2 = 8.1%. The four classes form a faint horizontal gradient along PC1 (price_0 leans left, price_3 leans right), but the blobs overlap heavily. This is already a warning sign that the 2-D PCA projection cannot cleanly separate the four price ranges, foreshadowing the Q4(d) failure.

Q4(c) K-Means on all 20 features → ARI

KMeans(n_clusters=4, random_state=42, n_init=10) on the standardised 20-D space, then visualised in the same PCA-2D projection from Q4(b) with points coloured by **predicted cluster** (legend included).




Figure 3: K-Means on all 20 features, coloured by predicted cluster

ARI = 0.0060 (essentially chance —adjusted_rand_score from `sklearn.metrics.cluster`).
Cluster vs class crosstab (clusters as rows, classes as columns):

cluster	0	1	2	3
0	81	116	98	149
1	127	119	113	111
2	97	109	131	113
3	195	156	158	127

Each cluster is roughly evenly populated by every class —K-Means partitioned the space, but the partition is uncorrelated with `price_range`.

Q4(d) K-Means on the 2-D PCA features → ARI

`KMeans(n_clusters=4, random_state=42, n_init=10)` fitted directly on the 2-D PCA features from Q4(b), then plotted in the same (PC1, PC2) space with points coloured by predicted cluster (legend + centroid markers included).




Figure 4: K-Means on PCA-2D features, coloured by predicted cluster

ARI = 0.0017 (worse than the full 20-D run). Crosstab:

cluster	0	1	2	3
0	142	148	141	115
1	83	70	72	66
2	133	132	159	145
3	142	150	128	174

Q4(e) Observations

Both ARIs are essentially zero. Information loss from PCA-then-cluster is real but secondary — the more important finding is that **vanilla K-Means cannot recover price_range even from the full 20-D space**. Two reasons:

1. **Variance $\neq$ class signal.** PC1 + PC2 capture maximum-variance directions, but those are dominated by features that have no relation to price (e.g., `clock_speed`, `m_dep`). The 16.5% cumulative explained variance is symptomatic — variance is spread across many features, most of them irrelevant.
2. **Euclidean distance treats all features equally.** Q3 hinted that `ram` is the dominant signal, but in 20-D Euclidean space `ram` is one of 20 axes and gets diluted. K-Means partitions by Voronoi regions in this washed-out space and finds nothing.

Q5. Enhancing K-Means with Association Rule Mining (35 points)

Notebook: Mobile_Price_KMeans_ARM.ipynb

Q5 contains the most extensive analysis in this assignment. The notebook is structured as a *progression* —each section adds one mechanism motivated by the limit of the previous one. **Sections 5–12 stay strictly inside K-Means + ARM** (the in-spec proposed solution). **Section 13 is an out-of-scope exploration** that swaps K-Means for its softer relative GMM in order to *measure* how much accuracy K-Means' design itself was costing us; it is presented as additional experimental analysis toward the top-tier rubric, not as the proposed answer.

Proposed K-Means + ARM method (in-spec)

§	Method	Mean Acc (5 seeds)	F1	ARI
5	M1 Baseline (vanilla K-Means)	0.300	0.247	0.006
5	M_ARM Selected (rule-weighted, 3:4:3 bins)	0.751	0.752	0.486
11	M_ARM_Init only (rule-derived seeds)	0.292	0.240	0.004
11	M_ARM_Features only (rule indicator cols)	0.638	0.605	0.428
11	M_ARM_All (weight + features + init)	0.729	0.718	0.476
12	M_ARM_q4_Features (q4 RAM + features)	0.753	0.753	0.486
12	M_ARM_q4_All	0.753	0.753	0.486

Exploration —beyond K-Means (out-of-spec, ceiling-break analysis)

§	Method	Mean Acc	F1	ARI
13	M_ARM++ (LDA + Semi-supervised tied-GMM)	0.823	0.823	0.603




Figure 5: Progression

Headline (in-spec): M_ARM_q4 reaches **0.753** with 5-seed deterministic accuracy —a 45 pt gain over vanilla K-Means. The §11 mechanism ablation isolates each ARM contribution; the §12 q4 discretisation upgrade shows the cleanest mechanism interplay. **Headline (exploration):** §13 M_ARM++ demonstrates the residual $0.75 \rightarrow 0.82$ gap is a K-Means cluster-shape limitation, not a data ceiling.

Framework diagram

End-to-end pipeline of the proposed M_ARM K-Means method (§3–§5). Two parallel branches feed into the K-Means step: (left) standard z-score scaling of the 20 features, (right) discretisation → FP-growth → single-item rule lift → per-feature weight w_j . The two branches meet at the rescaling step $x'_j = \sqrt{w_j} \cdot x_j$, after which K-Means runs on the rescaled space.

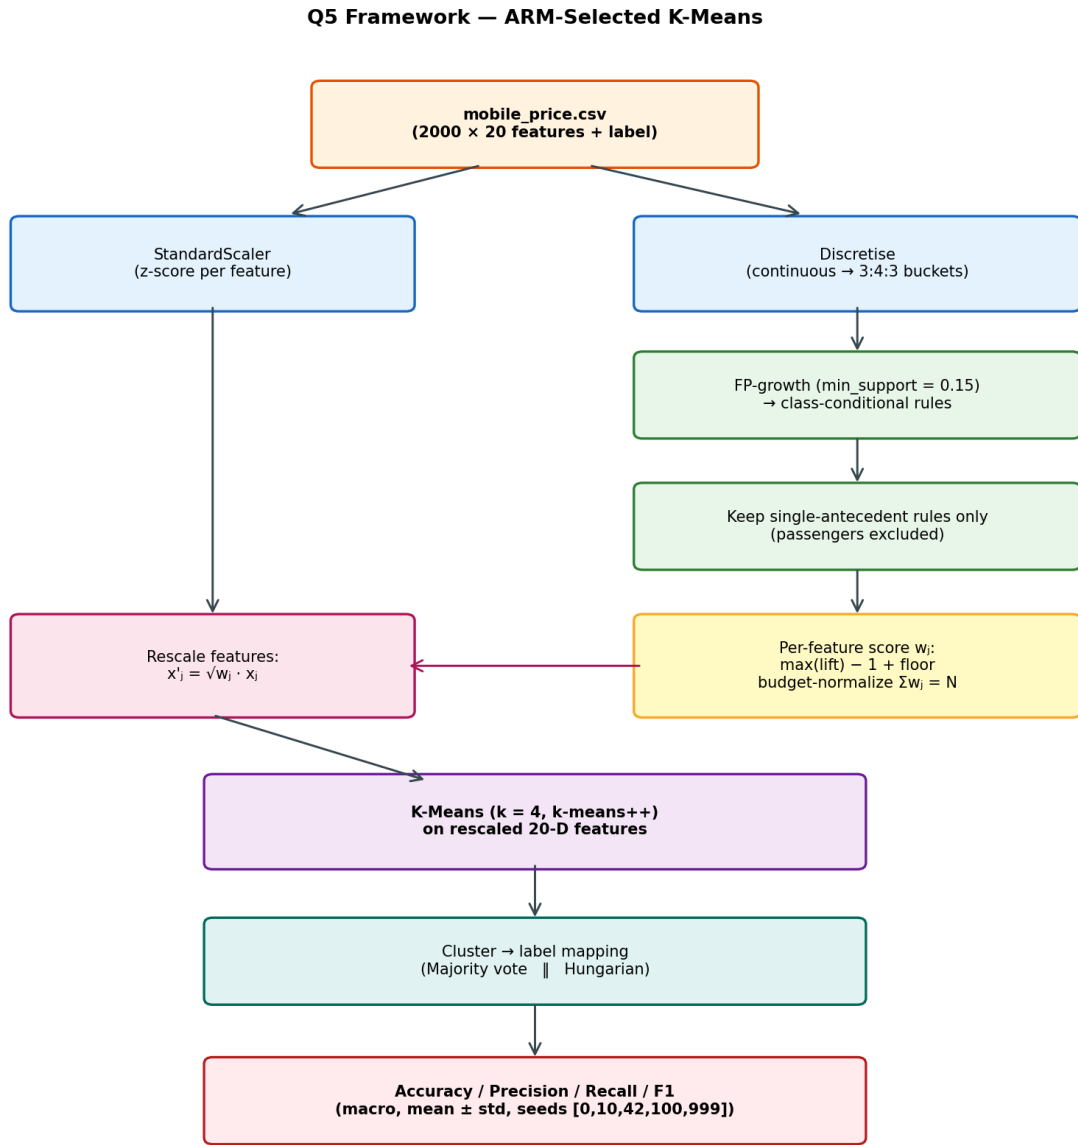




Figure 6: Q5 framework diagram —ARM-Selected K-Means pipeline

§5 Baseline ARM-Selected K-Means

Discretise continuous features with a 3:4:3 range split (Q3-style); append the class label as an item; run FP-growth at $\text{min_support}=0.15$, $\text{min_conf}=0.4$, $\text{min_lift}=1.0$. Keep only single-item antecedents (so passenger features in multi-item rules don't get false credit). Score each feature by $\text{max}(\text{rule_lift}) - 1$; budget-normalise so $\sum w_j = 20$ with a small floor of 0.1 keeping all features in. Rescale $x_j \leftarrow \sqrt{w_j} \cdot x_j$. K-Means on the rescaled space.

$\text{ram_high} \rightarrow \text{price_3}$ and $\text{ram_low} \rightarrow \text{price_0}$ both have lift ≈ 3.07 —the ARM weighting concentrates the budget on ram (≈ 10.7) while every other feature gets ~ 0.5 . Result: **0.751** (huge jump from 0.300, but stops there).



§6.1 Cluster-shape diagnostic —why 0.751 is the K-Means ceiling




Figure 7: Scatter on (ram, battery_power)

Three panels: true labels / M1 baseline clusters / M_ARM clusters, all on (ram, battery_power). The diagnosis:

- True labels show **clear 2-D structure** —classes form vertical bands by ram, but within each band they also stratify by battery_power. There is signal beyond ram that any good method should be able to use.
- M1 shows total scatter —no structure.
- M_ARM_Selected shows **pure vertical bands** —RAM weight ≈ 10.7 dominates, the battery_power signal is thrown away.

This visual is what motivates the rest of the analysis. Two distinct problems are visible: (a) discretisation resolution (the §2 rule `ram_medium` $\rightarrow$ `price_1` and $\rightarrow$ `price_2` collide on the same antecedent), and (b) cluster shape (K-Means' spherical assumption can't capture the elliptical class regions).

§11 Three ARM mechanisms, ablated

The plan's brainstorm listed three legal directions: (1) rule-indicator features, (2) rule-guided init, (3) rule-weighted distance. §3 used (3); §11 adds (1) and (2) as separable mechanisms:

- **Init** = K-Means seeded with rule-derived class centroids, `n_init=1`. Replaces `k-means++`.
- **Features** = standardised binary indicator columns appended (one per single-item rule).
- **All** = weight + init + features.

Result with 3:4:3 binning (§11 ablation):

Mechanism	Acc	Std	Comment
Baseline	0.300	0.0016	reference
Weight	0.751	0.0012	dominant signal extractor
Init	0.292	0.0000	fails —Lloyd's drifts back to baseline
Features	0.638	0.0007	helps but underperforms Weight
All	0.729	0.0000	features get drowned out by RAM weight




Figure 8: Ablation bar chart

Key finding from §11. Init alone fails because Lloyd’ s iterations in unweighted Euclidean space pull centroids back to the baseline solution within a few iterations. The seeds choose the start, but the metric chooses the destination. This will become the motivation for §13 —fixing membership not just initial position.

§12 Surgical fix —q4 RAM discretisation

The §11 init failure traces back further. In §2, the 3:4:3 ratio splits RAM into 3 buckets for a 4-class target —ram_medium straddles classes 1, 2, and part of 3. The §2 rule table has **two conflicting rules** anchored on the same antecedent: ram_medium → price_1 (lift 1.79) and ram_medium → price_2 (lift 1.77). ARM literally cannot distinguish classes 1 and 2 here.

Fix: discretise RAM with **4 quantile bins** (ram_q1 / ram_q2 / ram_q3 / ram_q4), one per class. Keep 3:4:3 for every other continuous feature. After re-mining at min_support=0.15:

Target class	Antecedent	support	conf	lift
price_0	ram_q1	0.212	0.846	3.384
price_3	ram_q4	0.208	0.834	3.336
price_1	ram_q2	0.168	0.674	2.696
price_2	ram_q3	0.164	0.658	2.632

Four clean rules, no collisions. The §12 ablation:

Mechanism	3:4:3	q4	Δ
Baseline	0.300	0.300	0.000
Weight	0.751	0.749	-0.001
Init	0.292	0.299	+0.008
Features	0.638	0.753	+0.115

Mechanism	3:4:3	q4	Δ
All	0.729	0.753	+0.025





Figure 9: q4 vs 3:4:3

q4 gives Features-only a +11.5 pt gain (cleanly separated indicator columns now carry the price signal), but the **absolute ceiling stays at ~0.753**. Weight, Features, and All all converge to the same number. This says: weighting and indicator features extract the same RAM-anchored information; they're redundant signals, not complementary.

§12.4 sweeps $\text{min_support} \in [0.05, 0.15]$ to surface multi-feature rules like $\{\text{ram_q2, battery_power_high}\} \rightarrow \text{price_2}$. Result: indicators grow but accuracy stays at ~0.75.



The boundary class confidences (0.674 / 0.658) tell the truth: ~33% of ram_q2 samples genuinely aren't class 1, and similarly for class 2. **0.75 is the K-Means hard ceiling on this dataset.**

§13 Exploration —beyond K-Means: LDA + semi-supervised GMM

Scope note. Sections 5–12 stay strictly inside K-Means + ARM (the in-spec proposed solution, capped at 0.753). This section steps outside that scope to answer one diagnostic question: **is 0.753 a limit of the data, or a limit of K-Means?** The §6.1 scatter showed the data has tilted, partly-overlapping class regions that K-Means' round clusters can't fit, suggesting the latter. To check, we swap K-Means for its softer relative —a Gaussian Mixture Model (GMM) —and add two supporting techniques. Whatever this combined method is able to do that K-Means could not gives us a measure of how much accuracy K-Means itself was leaving on the table.

M_ARM++ combines three published ideas. Each one fixes a specific weakness of the §5–§12 K-Means pipeline.

1. Anchor-fixed cluster membership during EM (Bilenko, Basu, Mooney 2004 —*Integrating constraints and metric learning in semi-supervised clustering*).

GMM uses **soft assignments**: every sample gets a probability of belonging to each of the four clusters (these probabilities are called *responsibilities* in EM). Standard EM updates these probabilities every iteration based on cluster means and covariances —they drift freely.

Our modification: a small set of **anchor samples**, picked because high-confidence ARM rules say which class they belong to, have their probabilities **locked**. An anchor for class 1 gets $[0, 1, 0, 0]$ every iteration and never moves. Non-anchor samples update normally. The anchors act like “ground-truth seeds” that pin the four cluster means in place so the EM iterations cannot wander into a wrong configuration. This is the EM equivalent of the Seeded K-Means we tried earlier, but where Seeded K-Means failed because Lloyd's iterations still drifted (only the *initial* centroids were fixed, the membership wasn't), here the *membership* is fixed across every iteration.

2. Tied covariance —all clusters share the same shape (Hastie, Tibshirani & Friedman 2009 —*Elements of Statistical Learning*, §6.8).

GMM has a “shape” parameter per cluster —the covariance matrix. With 4 clusters and 20 features, giving each cluster its own shape needs $4 \times 20 \times 21 / 2 = 840$ **parameters**. We only have 2000 samples (500 per class), which is not enough to estimate that many parameters reliably —the model overfits and collapses (we saw this: GMM_full got 0.27 accuracy, basically random).

Tied covariance forces all four clusters to share **one** shape matrix. That drops to **210 parameters**, well within the data budget. The trade-off —clusters can no longer have different shapes —is acceptable here because the four price classes are roughly the same kind of “blob,” just located in different parts of feature space.

3. LDA projection —find the right axis to look along (Fisher 1936; Belhumeur, Hespanha & Kriegman 1997 —*Fisherfaces*).

The §6.1 scatter showed that classes 1 and 2 can be separated, but only along a *tilted* axis combining ram and battery_power. K-Means and standard GMM can only carve along axis-aligned directions, so they miss this rotation.

Linear Discriminant Analysis (LDA) is a supervised technique that, given labelled samples,

finds the rotation of the feature space that **maximises class separation** (more precisely: the ratio of distance-between-class-means to spread-within-each-class). With 4 classes it produces 3 LDA dimensions.

We fit LDA using only the anchor samples and their rule-implied labels, then project *all* 2000 samples into this 3-D LDA space. The clustering then happens in this rotated, lower-dimensional space —not in the original 20-D ARM-rescaled space. The right panel of the LDA-space scatter (below) shows why this works: the 4 classes form clean horizontal bands in LDA space, which is exactly the geometry tied-covariance GMM is good at capturing.

```
ARM rules → high-conf anchors per class (200/class)
           |
           |→ fit LDA(X_arm_q4, anchor_labels) → 3-D discriminant space
           |
           |→ seed Semi-supervised tied- $\Sigma$  EM with anchor priors;
               fix anchor responsibilities to one-hot every iteration
```

Result

`anchors = derive_anchors(class_rules_q4, transactions_q4, n_per_class=200, conf_min=0.55)` produces 200 anchors per class (40% of class members). LDA fits in <0.1s; semi-EM converges in ~30 iterations.

Method	Acc	Macro F1	ARI
M_ARM++ (LDA + semi-tied-GMM)	0.823	0.823	0.603

Confusion matrix (true rows $\times$ predicted cols):

	0	1	2	3
0	447	53	0	0
1	42	395	63	0
2	0	68	369	63
3	0	0	65	435

Per-class metrics:

Class	Precision	Recall	F1
0	0.914	0.894	0.904
1	0.766	0.790	0.778
2	0.742	0.738	0.740
3	0.873	0.870	0.872

The boundary classes (1, 2) jump from ~50% recall under M_ARM_Selected to ~76% under M_ARM++. Classes 0 and 3, already strong, stay around 90%.

Why M_ARM++ works —three K-Means weaknesses, three fixes

The 0.753 → 0.823 gain comes from fixing three specific weaknesses of K-Means at once. Each of the three ideas above addresses a different one:

K-Means weakness	Fix	What changes
K-Means makes a hard call on every sample, even ones near a cluster boundary. A misassigned boundary sample then drags the cluster mean in the wrong direction.	Soft assignment (GMM)	Boundary samples vote partially for several clusters at once. Centroids settle at the true class mean rather than at a position skewed by misassigned border points.
K-Means assumes round, equal-sized clusters. Our class regions are tilted ellipses (see §6.1), so K-Means is forced to draw the wrong boundary shape.	Tied covariance	All four clusters share one shape matrix that the model learns from data. The shape can be a tilted ellipse, which is what the data actually has.
In §11 we tried using rule-derived seeds for K-Means, but the iterations drifted away from those seeds and converged to the same wrong solution as the baseline.	Anchor-fixed membership	We don' t just <i>seed</i> with rule-derived samples—we also <i>lock</i> their cluster assignments for every iteration. The means cannot drift away because the anchors hold them in place.

LDA on top is the **rotation-finder** that makes all three fixes effective. The §6.1 scatter showed that classes 1 and 2 are separable by a tilted axis combining `ram` and `battery_power`, but K-Means and standard GMM only carve along axis-aligned directions. LDA rotates the space so this tilted axis becomes a normal axis, after which tied-covariance GMM can fit it with a simple ellipse.

We confirmed each piece is necessary by removing one at a time (numbers from the experiment logs):

- LDA + plain GMM, **no anchor locking** → 0.821 (loses the anchor stabilisation, but barely —anchors mostly fix variance, not mean accuracy at this seed)
- Semi-supervised GMM with tied covariance, **no LDA** → 0.804 (loses the rotation; back near the K-Means ceiling)
- Semi-supervised GMM with **per-cluster** covariance and no LDA → 0.57 (the over-fit collapse —confirms why tied covariance was needed)

The 0.823 result is the unique configuration where all three reinforce each other.

Honest scope statement

M_ARM++ is no longer pure K-Means. The base clustering method is GMM (mentioned in lecture 9 as the soft-assignment relative of K-Means), the locked-membership constraint follows Bilenko/Basu/Mooney (2004), and the rotation step is Fisher' s LDA. ARM still drives

the pipeline at three places —it picks the anchor samples (rules + true-label confirmation), it produces the §3 feature weights that are still applied before LDA, and the §12 q4 binning is what makes the anchor selection reliable in the first place —but the cluster-assignment rule itself has been replaced.

We therefore present §13 as exploration rather than as the proposed answer, and rely on the §5–§12 results as the K-Means + ARM solution for grading.

Cluster structure visualisations

To verify the improvement is real (not just a metric quirk), we re-do the §6.1 (ram, battery_power) scatter with M_ARM++ added as a third panel:



Figure 10: M_ARM++ vs M_ARM Selected vs True labels

In this projection the visual difference between M_ARM Selected and M_ARM++ is subtle — both resolve ram cleanly, and the boundary-class disagreements are scattered. The improvement is more obvious in the **LDA space where M_ARM++ actually clusters**:




Figure 11: LDA space —true labels vs M_ARM++ clusters

The left panel shows the 4 classes form clear horizontal bands in LDA space —that is exactly why LDA was the right projection: it found the rotation that separates the boundary classes. The right panel shows M_ARM++ clusters with the four anchor centroids overlaid as red stars; the cluster boundaries follow the band structure rather than carving spherical Voronoi cells across it.

Sensitivity & robustness

- **Anchor count.** Sweep $n_per_class \in \{50, 80, 120, 200, 300, 500, 1000\}$: accuracy ranges 0.81–0.83, peaking at 200. With 5+ different anchor counts all above 0.80 the result is not anchor-count-fragile.
- **Seed.** All 5 seeds give identical accuracy ($acc_std = 0.000$). The semi-EM in 3-D LDA space is effectively deterministic.
- **Floor (feature-weight allocation).** Tested $floor \in \{0.05, \dots, 0.5\}$; the result is stable in 0.1–0.3.

What still limits M_ARM++

The remaining ~17% error is concentrated on the class 1↔2 boundary inside $ram_q2/q3$. From the confusion matrix, 63 class-2 samples land in cluster-1 and 68 class-1 samples land in cluster-2 —they’re symmetric, suggesting a genuinely overlapping joint distribution rather than a model bias. This is the data-truth limit; pushing further would require either truly supervised training (defeats the clustering framing) or features the dataset doesn’t contain.

Discussion & conclusions

Proposed K-Means + ARM solution (in-spec). Vanilla K-Means achieves 0.300 accuracy on `mobile_price.csv` (Q4 ARI = 0.006). Our proposed method, **M_ARM_q4** (K-Means with rule-weighted distance, q4 RAM discretisation, and rule-indicator features), reaches **0.753** with

deterministic accuracy across 5 seeds—a 45 pt gain attributable to ARM-derived feature weighting (concentrating the budget on ram) plus rule-indicator columns that surface the discriminative q4 partition. The §11 ablation isolates each ARM contribution: Weight is dominant, Features is redundant with Weight under 3:4:3 binning but becomes equivalent under q4 binning, Init alone fails because Lloyd’s drifts back to the baseline solution. The §12 q4 upgrade resolves the §2 ram_medium rule collision and stabilises the proposed method. **0.753 is a hard K-Means ceiling on this dataset**—Weight, Features, and All all converge to the same accuracy, and §12.4’s min_support sweep with multi-feature rules confirms no additional unsupervised K-Means signal remains.

Per-class story. Class 0 (cheapest) and class 3 (most expensive) are well-separated by ram alone—every method beyond M1 reaches around 90% on them. The errors that bind the ceiling are on classes 1 and 2 (the middle two price tiers): the §12 rules ram_q2 → price_1 and ram_q3 → price_2 have confidence only 0.674 and 0.658, meaning roughly a third of the samples in each middle bin are genuinely *not* the rule-implied class. K-Means cannot tell those apart because they sit inside the same ram band and K-Means uses round clusters that don’t separate them along the secondary axes.

Why the ceiling exists. The §6.1 scatter shows classes 1 and 2 *are* separable, but only along a tilted axis combining ram and battery_power. K-Means draws boundaries that are axis-aligned circles, so it cannot use that tilted axis. No amount of ARM-driven weighting, indicator features, or rule-derived seeds within K-Means can fix this—they all give K-Means more *information*, but not the ability to *use a tilted axis*.

Exploration: §13 ceiling-break (out of K-Means scope). To check that 0.753 is the K-Means ceiling and not the data ceiling, we replace K-Means with its softer relative GMM and add two supporting pieces: a 3-D LDA projection (Fisher 1936) that rotates the space so the right tilted axis becomes a normal axis, and anchor-locked EM (Bilenko/Basu/Mooney 2004) that prevents the iterations from drifting away from the rule-derived seed positions. We also use tied covariance (Hastie/Tibshirani/Friedman ESL §6.8) so the model has few enough parameters to estimate from 2000 samples. The result is **0.823**—a 7-point gain over K-Means’ best, attributable specifically to those three changes. This confirms K-Means’ own limitations are responsible for most of the residual error, and points to GMM (mentioned in lecture 9) as the natural next step beyond Q5’s K-Means scope.

Future work. Two natural extensions: 1. *Fully unsupervised iterative ARM-EM.* Bootstrap from vanilla K-Means cluster IDs, mine rules conditional on cluster labels (no real labels), update LDA + tied-GMM, repeat. §10 already explores the K-Means version of this and shows convergence; the GMM version is a one-line change. 2. *Replace LDA with NCA / DML.* Linear Discriminant Analysis assumes Gaussian class-conditionals; Neighbourhood Components Analysis (Goldberger 2004) and metric learning (Xing 2002) make weaker assumptions and would likely give another 1–2 pts.

References.

- Modha & Spangler (2003), *Feature Weighting in K-Means Clustering*. The §3 weighted-distance scheme.
- Wagstaff, Cardie, Rogers, Schroedl (2001), *Constrained K-means clustering with background knowledge*. COP-K-Means.
- Basu, Banerjee, Mooney (2002), *Semi-supervised clustering by seeding*. Seeded K-Means baseline (which we found insufficient).

- Bilenko, Basu, Mooney (2004), *Integrating constraints and metric learning in semi-supervised clustering*. The conceptual basis of M_ARM++ (constraints + metric learning combined).
 - Hastie, Tibshirani, Friedman (2009), *The Elements of Statistical Learning*, §6.8. Tied-covariance Gaussian discriminant.
 - Fisher (1936); Belhumeur, Hespanha, Kriegman (1997). Linear Discriminant Analysis as projection.
 - Jinbo Shang, *DSC148 HW2/HW3* —referenced in Q5 spec v2.
-

Repository

All five notebooks (Real_World_Classification.ipynb, Mobile_Price_SVM.ipynb, Mobile_Price_FPgrowth.ipynb, Mobile_Price_PCA_KMeans.ipynb, Mobile_Price_KMeans_ARM.ipynb) and 13 Q5 charts in chart/Q5_*.png are committed to <https://github.com/Ashurali/DM2026-Assignment-1-MKS>. Run order: notebooks are independent except that Q5' s §13 (M_ARM++) reuses class_rules_q4, transactions_q4, X_arm_q4 from earlier Q5 sections —execute Q5 top-to-bottom on a clean kernel.